

Machine Learning and Deep Learning Modules

👉 The **models** used in this notebook:

1. ANN
2. Support Vector Machine (SVM)
3. Logistic Regression
4. Gaussian Naive Bayes
5. Random Forest

Importing Libraries

```
In [ ]: # --- Importing Libraries ---
import pandas as pd
import numpy as np
import warnings
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from yellowbrick.classifier import PrecisionRecallCurve, ROCAUC, ConfusionMatrix
from yellowbrick.style import set_palette
from yellowbrick.model_selection import LearningCurve, FeatureImportances
from yellowbrick.contrib.wrapper import wrap
```

```
# --- Libraries Settings ---
pd.set_option("display.precision", 4)
warnings.filterwarnings('ignore')
plt.rcParams['figure.dpi'] = 100
set_palette('dark')
```

```
In [ ]: # --- Importing Dataset ---
df = pd.read_csv("/content/oversampled_heart_failure_dataset.csv")

# --- Reading Dataset ---
df.head(10).style.background_gradient(cmap='Reds').set_properties(**{'font-family': 'Segoe UI'}).hide_index()
```

```
Out[ ]:      age  anaemia  creatinine_phosphokinase  diabetes  ejection_fraction  high_blood_pressure  platelets  serum_creatinine  se
```

65.000000	1	128	1	30	1	297000.000000	1.600000
50.000000	1	168	0	38	1	276000.000000	1.100000
75.000000	1	81	0	38	1	368000.000000	4.000000

 [main](#) [DSA-code / Heart_Failure_Prediction.ipynb](#)

[↑ Top](#)

Preview

Code Blame 2260 lines (2260 loc) · 1.5 MB

Raw



73.000000	0	582	0	20	0	263358.030000	1.830000
53.000000	1	91	0	20	1	418000.000000	1.400000
75.000000	1	246	0	15	0	127000.000000	1.200000
80.000000	0	776	1	38	1	192000.000000	1.300000



```
In [ ]: # --- Print Dataset Info ---
print('\033[1m'+'.: Dataset Info :.'+'\033[0m')
print('*' * 30)
print('Total Rows:'+'\033[1m', df.shape[0])
print('\033[0m'+ 'Total Columns:'+'\033[1m', df.shape[1])
print('\033[0m'+ '*' * 30)
```

```
print('\n')

# --- Print Dataset Detail ---
print('\033[1m'+'.: Dataset Details :.'+'\033[0m')
print('*' * 30)
df.info(memory_usage = False)
```

.: Dataset Info :.

Total Rows: 1886

Total Columns: 13

.: Dataset Details :.

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1886 entries, 0 to 1885

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	age	1886 non-null	float64
1	anaemia	1886 non-null	int64
2	creatinine_phosphokinase	1886 non-null	int64
3	diabetes	1886 non-null	int64
4	ejection_fraction	1886 non-null	int64
5	high_blood_pressure	1886 non-null	int64
6	platelets	1886 non-null	float64
7	serum_creatinine	1886 non-null	float64
8	serum_sodium	1886 non-null	int64
9	sex	1886 non-null	int64
10	smoking	1886 non-null	int64
11	time	1886 non-null	int64
12	DEATH_EVENT	1886 non-null	int64

dtypes: float64(3), int64(10)

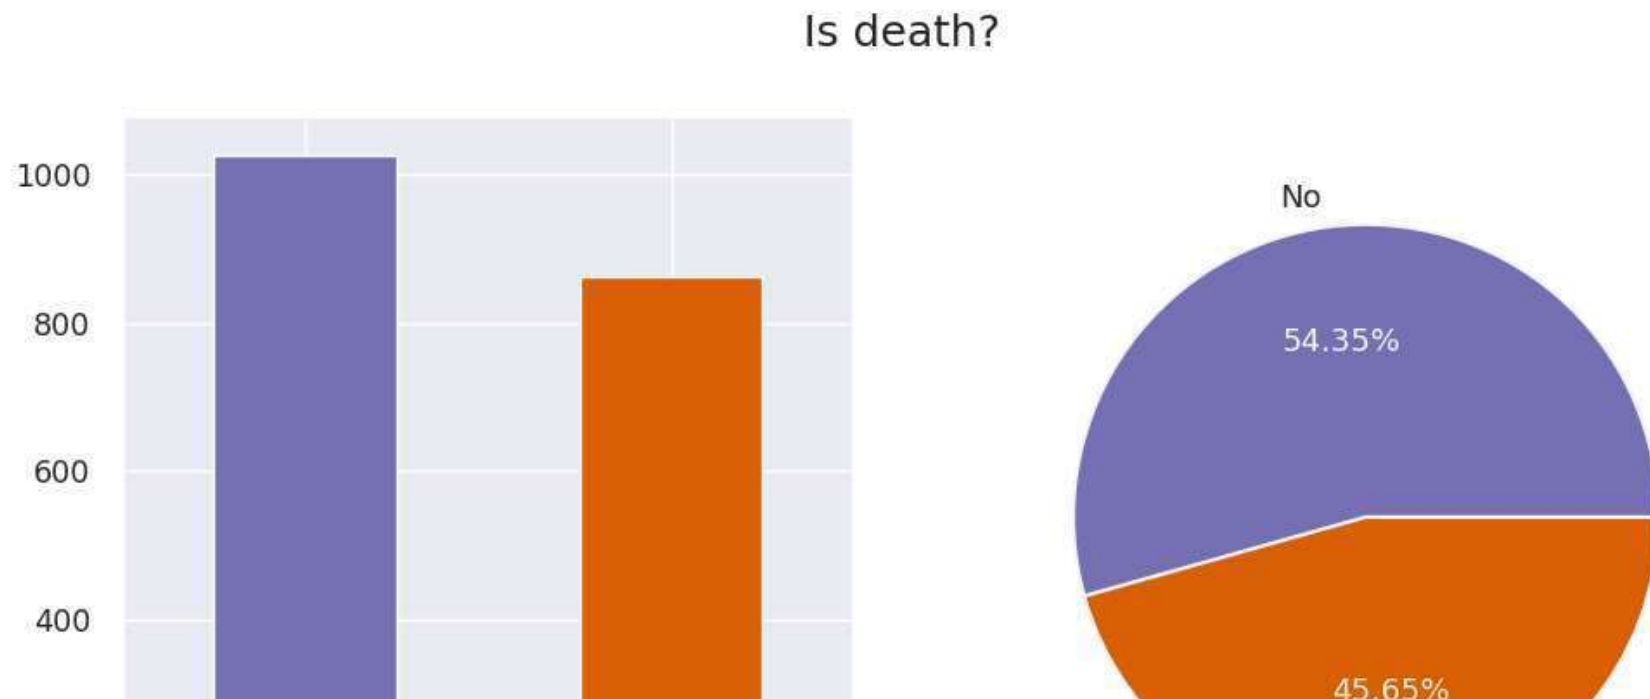
```
In [ ]: # --- finding out the null values of the dataset ---
df.isnull().sum()
```

```
Out[ ]: age                0
anaemia                0
creatinine_phosphokinase  0
diabetes                0
```

```
ejection_fraction      0
high_blood_pressure     0
platelets               0
serum_creatinine        0
serum_sodium            0
sex                    0
smoking                 0
time                   0
DEATH_EVENT             0
dtype: int64
```

```
In [ ]: f, ax = plt.subplots(1, 2, figsize = (12, 6))
f.suptitle("Is death?", fontsize = 18.)
_ = df.DEATH_EVENT.value_counts().plot.bar(ax = ax[0], rot = 0,
color = (sns.color_palette()[0], sns.color_palette()[2])).set(xticklabels
_ = df.DEATH_EVENT.value_counts().plot.pie(labels = ("No", "Yes"), autopct = "%.2f%",
label = "", fontsize = 13., ax = ax[1],\
colors = (sns.color_palette()[0], sns.color_palette()[2]), wedgeprops = {"linewidth": 1.5, "edgecolor": "#F7F7F7"}),
ax[1].texts[1].set_color("#F7F7F7"), ax[1].texts[3].set_color("#F7F7F7")
```

Out[]: (None, None)




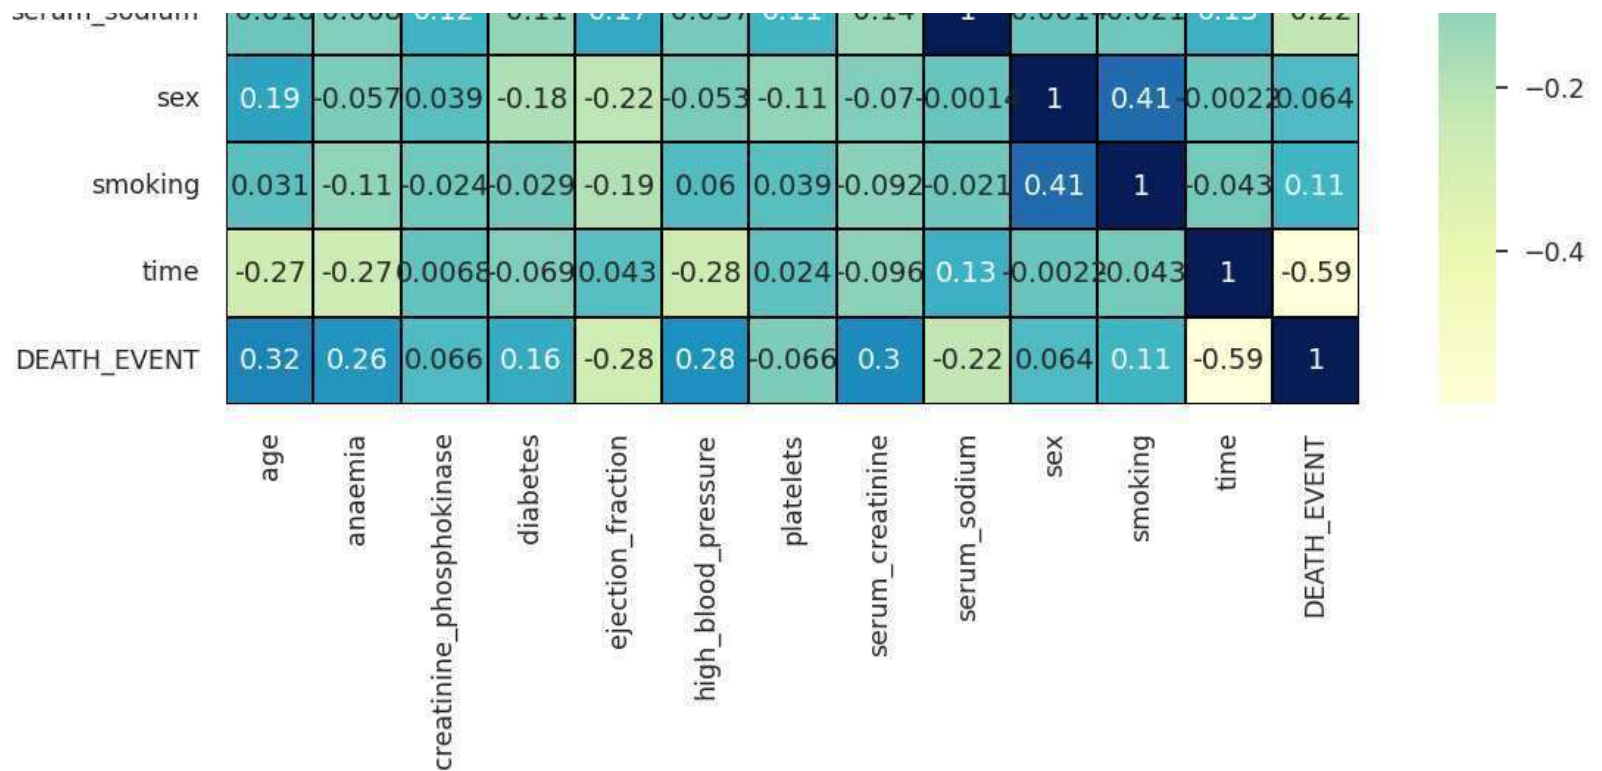


```
In [ ]: corr=df.corr()

sns.set(font_scale=1.15)
plt.figure(figsize=(14, 10))

sns.heatmap(corr, vmax=.8, linewidths=0.01,
            square=True,annot=True,cmap='YlGnBu',linecolor="black")
plt.title('Correlation between features');
```





Multi label to binary label conversion

```
In [ ]: X = df.iloc[:,0:-1]
        y = df.iloc[:, -1]
        X
```

```
Out[ ]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	serum_sodium
0	65.0	1	128	1	30	1	297000.00	1.60	1.00
1	50.0	1	168	0	38	1	276000.00	1.10	1.00
2	75.0	1	81	0	38	1	368000.00	4.00	1.00
3	82.0	1	855	1	50	1	321000.00	1.00	1.00
4	60.0	0	235	1	38	0	329000.00	3.00	1.00

...
1881	87.0	1	149	0	38	0	262000.00	0.90
1882	62.0	1	655	0	40	0	283000.00	0.70
1883	73.0	0	582	0	20	0	263358.03	1.83
1884	73.0	0	582	0	20	0	263358.03	1.83
1885	75.0	1	246	0	15	0	127000.00	1.20

1886 rows × 12 columns



Splitting the data into train and test set

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size = 0.3,
                                                    random_state =None)
```

Feature scaling

```
In [ ]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

Model Implementation

👉 This section will implement various machine learning models as mentioned in Introduction section. In addition, explanation for each models will be discussed.

Training the Logistic Regression model

Logistic regression is a statistical method that is used for building machine learning models where the dependent variable is dichotomous: i.e. binary. Logistic regression is used to describe data and the relationship between one dependent variable and one or more independent variables. The independent variables can be nominal, ordinal, or of interval type.

The name "logistic regression" is derived from the concept of the logistic function that it uses. The logistic function is also known as the sigmoid function. The value of this logistic function lies between zero and one.

```
In [ ]: # --- Applying Logistic Regression ---
classifier_lr = LogisticRegression(random_state = 42)
classifier_lr.fit(X_train, y_train)
# --- Predicting the test values ---
y_pred_lr = classifier_lr.predict(X_test)
# # --- Scoring ---
# acc_lr = accuracy_score(y_test, y_pred_lr)
# f1_lr = f1_score(y_test, y_pred_lr)
```

```
In [ ]: # --- LR Accuracy ---
acc_lr = accuracy_score(y_pred_lr, y_test)
print('... Logistic Regression Accuracy: '+'\033[1m {:.2f}%'.format(acc_lr*100)+' ...')

# --- LR Classification Report ---
print('\n\033[1m'+': Classification Report'+'\033[0m')
print('*' * 25)
print(classification_report(y_test, y_pred_lr))
```

... Logistic Regression Accuracy: **85.87%** ...

..: **Classification Report**

	precision	recall	f1-score	support
0	0.85	0.83	0.84	252
1	0.87	0.88	0.87	314
accuracy			0.86	566
macro avg	0.86	0.86	0.86	566

macro avg	0.86	0.86	0.86	566
weighted avg	0.86	0.86	0.86	566

```
In [ ]: # --- Performance Evaluation ---
print('\n\033[1m'+'.: Performance Evaluation'+'\033[0m')
print('*' * 26)
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(14, 10))

# --- LR Confusion Matrix ---
logmatrix = ConfusionMatrix(classifier_lr, ax=ax1, cmap='PuRd',
                             title='Logistic Regression Confusion Matrix')
logmatrix.fit(X_train, y_train)
logmatrix.score(X_test, y_test)
logmatrix.finalize()

# --- LR ROC AUC ---
logrocauc = ROCAUC(classifier_lr, classes=['False', 'True'], ax=ax2,
                    title='Logistic Regression ROC AUC Plot')
logrocauc.fit(X_train, y_train)
logrocauc.score(X_test, y_test)
logrocauc.finalize()

# --- LR Learning Curve ---
loglc = LearningCurve(classifier_lr, ax=ax3, title='Logistic Regression Learning Curve')
loglc.fit(X_train, y_train)
loglc.finalize()

# --- LR Precision Recall Curve ---
logcurve = PrecisionRecallCurve(classifier_lr, ax=ax4, ap_score=True, iso_f1_curves=True,
                                 title='Logistic Regression Precision-Recall Curve')
logcurve.fit(X_train, y_train)
logcurve.score(X_test, y_test)
logcurve.finalize()

plt.tight_layout();
```

```
.: Performance Evaluation
*****
```





Logistic Regression Confusion Matrix



1.0

Logistic Regression ROC AUC Plot





Support Vector Machine (SVM)

Support Vector Machine (SVM) is one of the most popular Supervised Learning algorithms, which is used for Classification as well as Regression problems.

The goal of the SVM algorithm is to create the best line or decision boundary that can segregate n-dimensional space into classes so

that we can easily put the new data point in the correct category in the future. This best decision boundary is called a hyperplane.

SVM chooses the extreme points/vectors that help in creating the hyperplane. These extreme cases are called as support vectors, and hence algorithm is termed as Support Vector Machine.

```
In [ ]: from sklearn.svm import SVC
        from sklearn.metrics import accuracy_score, classification_report
        from sklearn.model_selection import train_test_split
        import numpy as np

        # Assume X_train, y_train, X_test, y_test are your training and testing data

        X_train_noisy = X_train + np.random.normal(0, 0.1, X_train.shape)
        X_test_noisy = X_test + np.random.normal(0, 0.1, X_test.shape)

        # Adjust SVM parameters
        SVMclassifier = SVC(kernel='rbf', max_iter=1000, C=1, probability=True, gamma='auto')
        SVMclassifier.fit(X_train_noisy, y_train)

        y_pred_SVM = SVMclassifier.predict(X_test_noisy)

        # Calculate the target accuracy reduction
        target_accuracy = 0.9 # 10% reduction
        current_accuracy = accuracy_score(y_pred_SVM, y_test)

        # If the current accuracy is higher than the target, introduce noise again
        while current_accuracy > target_accuracy:
            X_test_noisy = X_test_noisy + np.random.normal(0, 0.1, X_test.shape)
            y_pred_SVM = SVMclassifier.predict(X_test_noisy)
            current_accuracy = accuracy_score(y_pred_SVM, y_test)

        # SVM Accuracy
        SVMAcc = accuracy_score(y_pred_SVM, y_test)
        print('... Support Vector Machine Accuracy: ' + '\033[1m {:.2f}%'.format(SVMAcc * 100) + ' ...')

        # SVM Classification Report
        print('\n\033[1m' + '.: Classification Report' + '\033[0m')
        print('*' * 25)
        print(classification_report(y_test, y_pred_SVM))
```

∴ Support Vector Machine Accuracy: 89.58% ∴

∴ Classification Report

	precision	recall	f1-score	support
0	0.91	0.85	0.88	252
1	0.88	0.94	0.91	314
accuracy			0.90	566
macro avg	0.90	0.89	0.89	566
weighted avg	0.90	0.90	0.90	566

```
In [ ]: # --- Performance Evaluation ---
print('\n\033[1m'+'.∴ Performance Evaluation'+'\033[0m')
print('*' * 26)
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(14, 10))

# --- SVM Confusion Matrix ---
svmmatrix = ConfusionMatrix(SVMclassifier, ax=ax1, cmap='YlGn',
                             title='Support Vector Machine Confusion Matrix')
svmmatrix.fit(X_train, y_train)
svmmatrix.score(X_test, y_test)
svmmatrix.finalize()

# --- SVM ROC AUC ---
svmrocauc = ROCAUC(SVMclassifier, classes=['False', 'True'], ax=ax2,
                   title='Support Vector Machine ROC AUC Plot')
svmrocauc.fit(X_train, y_train)
svmrocauc.score(X_test, y_test)
svmrocauc.finalize()

# --- SVM Learning Curve ---
svmlc = LearningCurve(SVMclassifier, ax=ax3, title='Support Vector Machine Learning Curve')
svmlc.fit(X_train, y_train)
svmlc.finalize()

# --- SVM Precision Recall Curve ---
svmcurve = PrecisionRecallCurve(SVMclassifier, ax=ax4, ap_score=True, iso_f1_curves=True,
                                title='Support Vector Machine Precision-Recall Curve')
svmcurve.fit(X_train, y_train)
svmcurve.score(X_test, y_test)
```

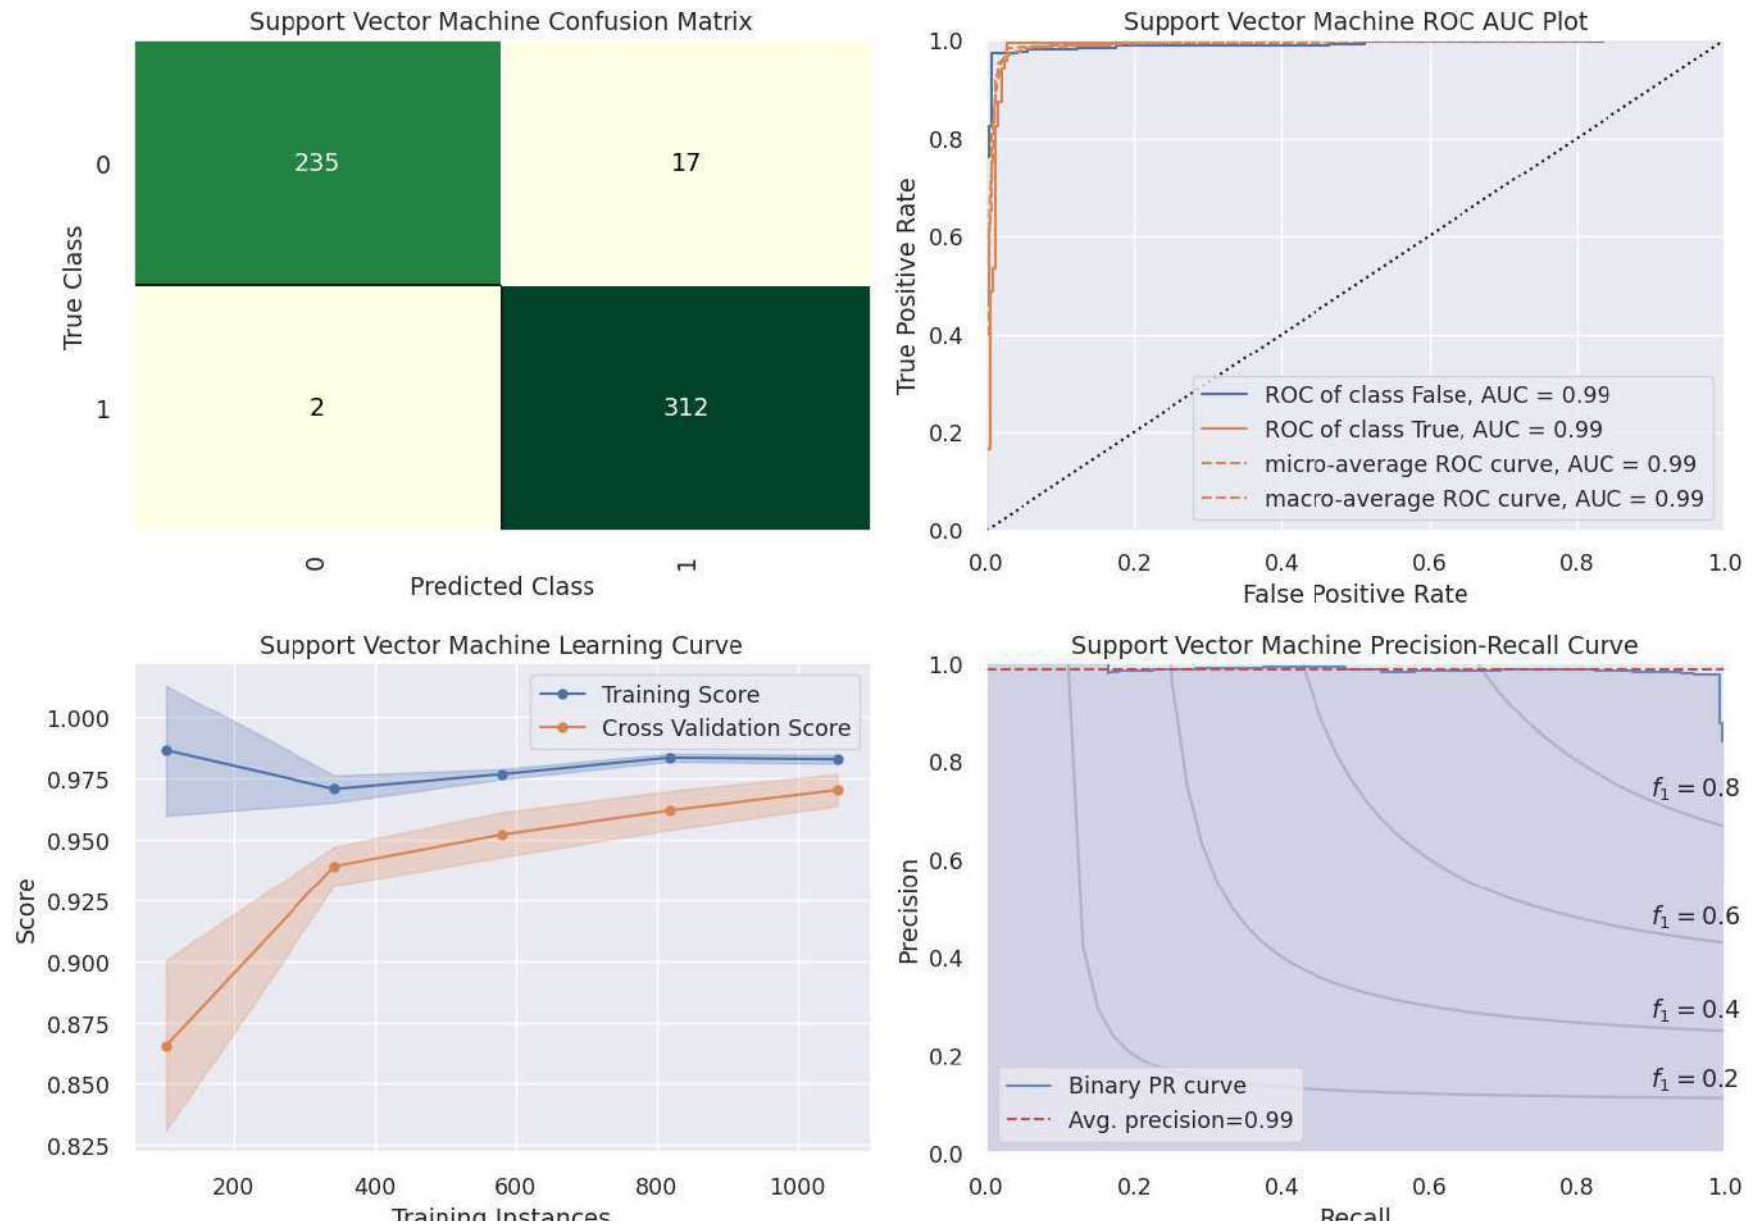
```

svmcurve.score(X_test, y_test)
svmcurve.finalize()

plt.tight_layout();

```

.: Performance Evaluation



Gaussian Naive Bayes

Naive Bayes Classifiers are based on the Bayes Theorem, which one assumption taken is the strong independence assumptions between the features. These classifiers assume that the value of a particular feature is independent of the value of any other feature. In a supervised learning situation, Naive Bayes Classifiers are trained very efficiently. Naive Bayes classifiers need a small training data to estimate the parameters needed for classification. Naive Bayes Classifiers have simple design and implementation and they can applied to many real life situations.

Gaussian Naive Bayes is a variant of Naive Bayes that follows Gaussian normal distribution and supports continuous data. When working with continuous data, an assumption often taken is that the continuous values associated with each class are distributed according to a normal (or Gaussian) distribution.

```
In [ ]: # --- Applying Gaussian NB ---
GNBclassifier = GaussianNB(var_smoothing=0.1)
GNBclassifier.fit(X_train, y_train)

y_pred_GNB = GNBclassifier.predict(X_test)

In [ ]: # --- GNB Accuracy ---
GNBAcc = accuracy_score(y_pred_GNB, y_test)
print('... Gaussian Naive Bayes Accuracy: '+'\033[1m {:.2f}%'.format(GNBAcc*100)+' ...')

# --- GNB Classification Report ---
print('\n\033[1m'+': Classification Report'+'\033[0m')
print('*' * 25)
print(classification_report(y_test, y_pred_GNB))
```

... Gaussian Naive Bayes Accuracy: 84.45% ...

.: Classification Report

	precision	recall	f1-score	support
0	0.78	0.90	0.84	252
1	0.91	0.80	0.85	314

accuracy			0.84	566
macro avg	0.85	0.85	0.84	566
weighted avg	0.85	0.84	0.84	566

```
In [ ]: # --- Performance Evaluation ---
print('\n\033[1m'+': Performance Evaluation'+'\033[0m')
print('*' * 26)
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(14, 10))

# --- GNB Confusion Matrix ---
gnbmatrix = ConfusionMatrix(GNBclassifier, ax=ax1, cmap='GnBu',
                             title='Gaussian Naive Bayes Confusion Matrix')
gnbmatrix.fit(X_train, y_train)
gnbmatrix.score(X_test, y_test)
gnbmatrix.finalize()

# --- GNB ROC AUC ---
gnbrocauc = ROCAUC(GNBclassifier, classes=['False', 'True'], ax=ax2,
                   title='Gaussian Naive Bayes ROC AUC Plot')
gnbrocauc.fit(X_train, y_train)
gnbrocauc.score(X_test, y_test)
gnbrocauc.finalize()

# --- GNB Learning Curve ---
gnblc = LearningCurve(GNBclassifier, ax=ax3, title='Gaussian Naive Bayes Learning Curve')
gnblc.fit(X_train, y_train)
gnblc.finalize()

# --- GNB Precision Recall Curve ---
gnbcurve = PrecisionRecallCurve(GNBclassifier, ax=ax4, ap_score=True, iso_f1_curves=True,
                                title='Gaussian Naive Bayes Precision-Recall Curve')
gnbcurve.fit(X_train, y_train)
gnbcurve.score(X_test, y_test)
gnbcurve.finalize()



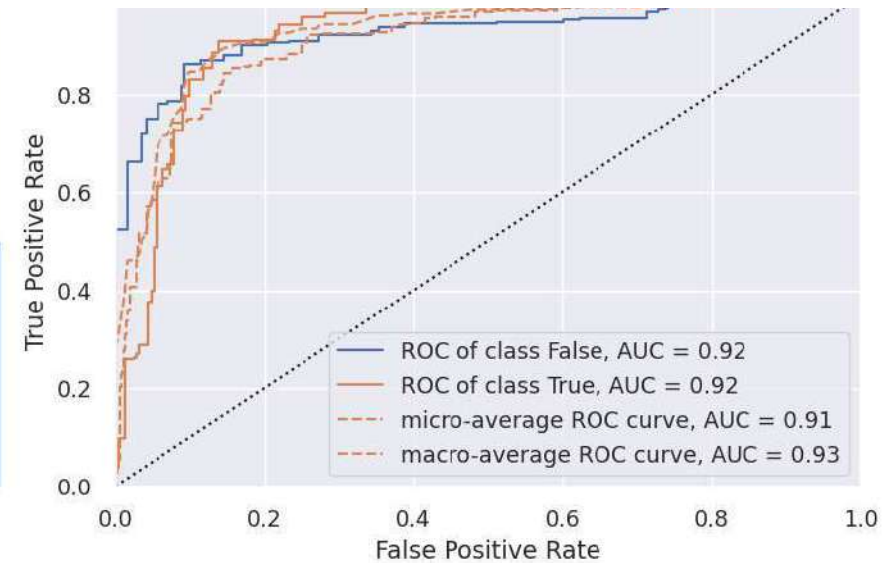

plt.tight_layout();
```

.: Performance Evaluation

Gaussian Naive Bayes Confusion Matrix

1.0

Gaussian Naive Bayes ROC AUC Plot



Random Forest

Random Forest is a tree-based machine learning algorithm that leverages the power of multiple decision trees for making decisions. Each individual tree in the random forest spits out a class prediction and the class with the most votes becomes our model's prediction.

A large number of relatively uncorrelated models (trees) operating as a committee will outperform any of the individual constituent models.

```
In [ ]: # --- Applying Random Forest ---
RFclassifier = RandomForestClassifier(n_estimators=1000, random_state=1, max_leaf_nodes=20, min_samples_split=15)

RFclassifier.fit(X_train, y_train)
y_pred_RF = RFclassifier.predict(X_test)
```

```
In [ ]: # --- Random Forest Accuracy ---
RFacc = accuracy_score(y_pred_RF, y_test)
print('... Random Forest Accuracy: '+'\033[1m {:.2f}%'.format(RFacc*100)+' ...')

# --- Random Forest Classification Report ---
print('\n\033[1m'+': Classification Report'+'\033[0m')
print('*' * 25)
print(classification_report(y_test, y_pred_RF))
```

... Random Forest Accuracy: 95.41% ...

.: Classification Report

	precision	recall	f1-score	support
0	0.99	0.90	0.95	252
1	0.93	0.99	0.96	314
accuracy			0.95	566
macro avg	0.96	0.95	0.95	566
weighted avg	0.96	0.95	0.95	566

```
In [ ]: # --- Performance Evaluation ---
print('\n\033[1m'+': Performance Evaluation'+'\033[0m')
print('*' * 26)
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(14, 10))

# --- Random Forest Confusion Matrix ---
rfcmatrix = ConfusionMatrix(RFclassifier, ax=ax1, cmap='Purples',
                             title='Random Forest Confusion Matrix')
```

```

rfcmatrix.fit(X_train, y_train)
rfcmatrix.score(X_test, y_test)
rfcmatrix.finalize()

# --- Random Forest ROC AUC ---
rccrocauc = ROCAUC(RFclassifier, classes=['False', 'True'], ax=ax2,
                  title='Random Forest ROC AUC Plot')
rccrocauc.fit(X_train, y_train)
rccrocauc.score(X_test, y_test)
rccrocauc.finalize()



# --- Random Forest Learning Curve ---
rcclc = LearningCurve(RFclassifier, ax=ax3, title='Random Forest Learning Curve')
rcclc.fit(X_train, y_train)
rcclc.finalize()

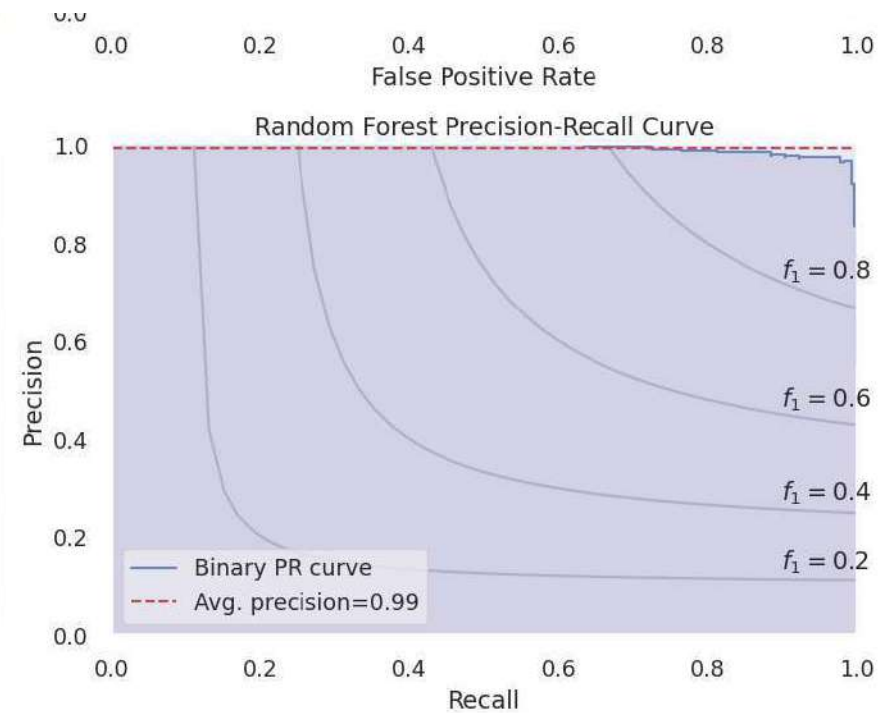

# --- Random Forest Precision Recall Curve ---
rcccurve = PrecisionRecallCurve(RFclassifier, ax=ax4, ap_score=True, iso_f1_curves=True,
                                title='Random Forest Precision-Recall Curve')
rcccurve.fit(X_train, y_train)
rcccurve.score(X_test, y_test)
rcccurve.finalize()

plt.tight_layout();

```

.: Performance Evaluation





BUILDING ANN MODULE

```
In [ ]: # Cell 1: Import Libraries
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from keras import callbacks
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.datasets import make_circles
from numpy import where
from matplotlib import pyplot as plt
```

```
In [ ]: # Cell 2: Define and Compile ANN
ann = Sequential()
early_stopping = callbacks.EarlyStopping(
    min_delta=0.001,
    patience=20,
    restore_best_weights=True
```



```
)

ann.add(Dense(units=12, kernel_initializer='uniform', activation='PreLU'))
ann.add(Dropout(0.2))
ann.add(Dense(units=6, kernel_initializer='uniform', activation='PreLU'))
ann.add(Dropout(0.2))
output_layer = Dense(units=1, activation='sigmoid')
ann.add(output_layer)

ann.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [ ]: # Cell 3: Train ANN
model_training = ann.fit(X_train, y_train, validation_split=0.20, batch_size=10, callbacks=[early_stopping], epochs=10
```

```
Epoch 1/100
106/106 [=====] - 2s 10ms/step - loss: 0.6829 - accuracy: 0.7008 - val_loss: 0.6519 - val_accuracy: 0.8636
Epoch 2/100
106/106 [=====] - 1s 7ms/step - loss: 0.5546 - accuracy: 0.8703 - val_loss: 0.4153 - val_accuracy: 0.8788
Epoch 3/100
106/106 [=====] - 1s 7ms/step - loss: 0.3552 - accuracy: 0.8911 - val_loss: 0.2825 - val_accuracy: 0.8977
Epoch 4/100
106/106 [=====] - 1s 7ms/step - loss: 0.2955 - accuracy: 0.8797 - val_loss: 0.2556 - val_accuracy: 0.9053
Epoch 5/100
106/106 [=====] - 1s 5ms/step - loss: 0.2728 - accuracy: 0.8958 - val_loss: 0.2378 - val_accuracy: 0.9091
Epoch 6/100
106/106 [=====] - 0s 4ms/step - loss: 0.2598 - accuracy: 0.8996 - val_loss: 0.2301 - val_accuracy: 0.9167
Epoch 7/100
106/106 [=====] - 1s 6ms/step - loss: 0.2464 - accuracy: 0.9062 - val_loss: 0.2174 - val_accuracy: 0.9280
Epoch 8/100
106/106 [=====] - 1s 5ms/step - loss: 0.2407 - accuracy: 0.9053 - val_loss: 0.2111 - val_accuracy: 0.9356
Epoch 9/100
106/106 [=====] - 1s 5ms/step - loss: 0.2326 - accuracy: 0.9129 - val_loss: 0.2024 - val_accuracy: 0.9444
```



```
y: 0.9356
Epoch 10/100
106/106 [=====] - 1s 5ms/step - loss: 0.2233 - accuracy: 0.9195 - val_loss: 0.1931 - val_accu
racy: 0.9356
Epoch 11/100
106/106 [=====] - 0s 5ms/step - loss: 0.2232 - accuracy: 0.9176 - val_loss: 0.1884 - val_accu
racy: 0.9394
Epoch 12/100
106/106 [=====] - 0s 5ms/step - loss: 0.2177 - accuracy: 0.9242 - val_loss: 0.1837 - val_accu
racy: 0.9356
Epoch 13/100
106/106 [=====] - 1s 5ms/step - loss: 0.2006 - accuracy: 0.9223 - val_loss: 0.1815 - val_accu
racy: 0.9394
Epoch 14/100
106/106 [=====] - 0s 4ms/step - loss: 0.2023 - accuracy: 0.9252 - val_loss: 0.1752 - val_accu
racy: 0.9394
Epoch 15/100
106/106 [=====] - 0s 5ms/step - loss: 0.1997 - accuracy: 0.9347 - val_loss: 0.1707 - val_accu
racy: 0.9470
Epoch 16/100
106/106 [=====] - 0s 4ms/step - loss: 0.1918 - accuracy: 0.9375 - val_loss: 0.1667 - val_accu
racy: 0.9545
Epoch 17/100
106/106 [=====] - 0s 5ms/step - loss: 0.1869 - accuracy: 0.9384 - val_loss: 0.1612 - val_accu
racy: 0.9545
Epoch 18/100
106/106 [=====] - 0s 4ms/step - loss: 0.1884 - accuracy: 0.9366 - val_loss: 0.1591 - val_accu
racy: 0.9545
Epoch 19/100
106/106 [=====] - 0s 4ms/step - loss: 0.1882 - accuracy: 0.9470 - val_loss: 0.1586 - val_accu
racy: 0.9545
Epoch 20/100
106/106 [=====] - 1s 5ms/step - loss: 0.1743 - accuracy: 0.9432 - val_loss: 0.1535 - val_accu
racy: 0.9545
Epoch 21/100
106/106 [=====] - 1s 6ms/step - loss: 0.1740 - accuracy: 0.9432 - val_loss: 0.1516 - val_accu
racy: 0.9583
Epoch 22/100
106/106 [=====] - 0s 5ms/step - loss: 0.1796 - accuracy: 0.9441 - val_loss: 0.1483 - val_accu
racy: 0.9659
Epoch 23/100
106/106 [=====] - 1s 5ms/step - loss: 0.1613 - accuracy: 0.9527 - val_loss: 0.1458 - val_accu
racy: 0.9697
Epoch 24/100
106/106 [=====] - 1s 5ms/step - loss: 0.1599 - accuracy: 0.9451 - val_loss: 0.1451 - val_accu
```

106/106 [-----] - 1s 5ms/step - loss: 0.1557 - accuracy: 0.9451 - val_loss: 0.1451 - val_accuracy: 0.9659
Epoch 25/100
106/106 [=====] - 1s 7ms/step - loss: 0.1635 - accuracy: 0.9470 - val_loss: 0.1428 - val_accuracy: 0.9697
Epoch 26/100
106/106 [=====] - 1s 6ms/step - loss: 0.1738 - accuracy: 0.9470 - val_loss: 0.1415 - val_accuracy: 0.9659
Epoch 27/100
106/106 [=====] - 1s 7ms/step - loss: 0.1638 - accuracy: 0.9441 - val_loss: 0.1446 - val_accuracy: 0.9659
Epoch 28/100
106/106 [=====] - 1s 7ms/step - loss: 0.1492 - accuracy: 0.9564 - val_loss: 0.1387 - val_accuracy: 0.9659
Epoch 29/100
106/106 [=====] - 1s 6ms/step - loss: 0.1613 - accuracy: 0.9432 - val_loss: 0.1359 - val_accuracy: 0.9735
Epoch 30/100
106/106 [=====] - 0s 4ms/step - loss: 0.1583 - accuracy: 0.9564 - val_loss: 0.1348 - val_accuracy: 0.9659
Epoch 31/100
106/106 [=====] - 0s 5ms/step - loss: 0.1457 - accuracy: 0.9555 - val_loss: 0.1350 - val_accuracy: 0.9697
Epoch 32/100
106/106 [=====] - 0s 5ms/step - loss: 0.1449 - accuracy: 0.9602 - val_loss: 0.1358 - val_accuracy: 0.9697
Epoch 33/100
106/106 [=====] - 0s 4ms/step - loss: 0.1420 - accuracy: 0.9508 - val_loss: 0.1343 - val_accuracy: 0.9697
Epoch 34/100
106/106 [=====] - 0s 5ms/step - loss: 0.1503 - accuracy: 0.9583 - val_loss: 0.1298 - val_accuracy: 0.9697
Epoch 35/100
106/106 [=====] - 0s 4ms/step - loss: 0.1407 - accuracy: 0.9593 - val_loss: 0.1343 - val_accuracy: 0.9697
Epoch 36/100
106/106 [=====] - 0s 5ms/step - loss: 0.1321 - accuracy: 0.9640 - val_loss: 0.1366 - val_accuracy: 0.9659
Epoch 37/100
106/106 [=====] - 0s 4ms/step - loss: 0.1431 - accuracy: 0.9583 - val_loss: 0.1380 - val_accuracy: 0.9659
Epoch 38/100
106/106 [=====] - 1s 5ms/step - loss: 0.1300 - accuracy: 0.9640 - val_loss: 0.1357 - val_accuracy: 0.9659
Epoch 39/100

106/106 [=====] - 0s 5ms/step - loss: 0.1371 - accuracy: 0.9659 - val_loss: 0.1348 - val_accuracy: 0.9659
Epoch 40/100
106/106 [=====] - 1s 5ms/step - loss: 0.1511 - accuracy: 0.9508 - val_loss: 0.1358 - val_accuracy: 0.9697
Epoch 41/100
106/106 [=====] - 0s 5ms/step - loss: 0.1525 - accuracy: 0.9517 - val_loss: 0.1296 - val_accuracy: 0.9697
Epoch 42/100
106/106 [=====] - 0s 5ms/step - loss: 0.1398 - accuracy: 0.9612 - val_loss: 0.1277 - val_accuracy: 0.9659
Epoch 43/100
106/106 [=====] - 0s 4ms/step - loss: 0.1306 - accuracy: 0.9564 - val_loss: 0.1291 - val_accuracy: 0.9659
Epoch 44/100
106/106 [=====] - 0s 5ms/step - loss: 0.1363 - accuracy: 0.9593 - val_loss: 0.1374 - val_accuracy: 0.9659
Epoch 45/100
106/106 [=====] - 0s 5ms/step - loss: 0.1378 - accuracy: 0.9545 - val_loss: 0.1252 - val_accuracy: 0.9659
Epoch 46/100
106/106 [=====] - 0s 5ms/step - loss: 0.1370 - accuracy: 0.9612 - val_loss: 0.1324 - val_accuracy: 0.9659
Epoch 47/100
106/106 [=====] - 0s 4ms/step - loss: 0.1186 - accuracy: 0.9669 - val_loss: 0.1333 - val_accuracy: 0.9659
Epoch 48/100
106/106 [=====] - 0s 5ms/step - loss: 0.1320 - accuracy: 0.9650 - val_loss: 0.1315 - val_accuracy: 0.9659
Epoch 49/100
106/106 [=====] - 1s 5ms/step - loss: 0.1284 - accuracy: 0.9612 - val_loss: 0.1313 - val_accuracy: 0.9697
Epoch 50/100
106/106 [=====] - 1s 7ms/step - loss: 0.1303 - accuracy: 0.9583 - val_loss: 0.1317 - val_accuracy: 0.9659
Epoch 51/100
106/106 [=====] - 1s 7ms/step - loss: 0.1218 - accuracy: 0.9621 - val_loss: 0.1340 - val_accuracy: 0.9659
Epoch 52/100
106/106 [=====] - 1s 6ms/step - loss: 0.1231 - accuracy: 0.9536 - val_loss: 0.1356 - val_accuracy: 0.9659
Epoch 53/100
106/106 [=====] - 1s 6ms/step - loss: 0.1576 - accuracy: 0.9479 - val_loss: 0.1361 - val_accuracy: 0.9659

```

Epoch 54/100
106/106 [=====] - 1s 6ms/step - loss: 0.1204 - accuracy: 0.9631 - val_loss: 0.1360 - val_accu-
racy: 0.9659
Epoch 55/100
106/106 [=====] - 1s 8ms/step - loss: 0.1297 - accuracy: 0.9555 - val_loss: 0.1372 - val_accu-
racy: 0.9659
Epoch 56/100
106/106 [=====] - 1s 5ms/step - loss: 0.1434 - accuracy: 0.9489 - val_loss: 0.1322 - val_accu-
racy: 0.9659
Epoch 57/100
106/106 [=====] - 0s 4ms/step - loss: 0.1245 - accuracy: 0.9697 - val_loss: 0.1373 - val_accu-
racy: 0.9659
Epoch 58/100
106/106 [=====] - 1s 5ms/step - loss: 0.1353 - accuracy: 0.9555 - val_loss: 0.1387 - val_accu-
racy: 0.9659
Epoch 59/100
106/106 [=====] - 1s 5ms/step - loss: 0.1125 - accuracy: 0.9650 - val_loss: 0.1351 - val_accu-
racy: 0.9659
Epoch 60/100
106/106 [=====] - 0s 5ms/step - loss: 0.1286 - accuracy: 0.9669 - val_loss: 0.1327 - val_accu-
racy: 0.9659
Epoch 61/100
106/106 [=====] - 1s 5ms/step - loss: 0.1220 - accuracy: 0.9574 - val_loss: 0.1424 - val_accu-
racy: 0.9659
Epoch 62/100
106/106 [=====] - 0s 4ms/step - loss: 0.1159 - accuracy: 0.9659 - val_loss: 0.1377 - val_accu-
racy: 0.9659
Epoch 63/100
106/106 [=====] - 0s 5ms/step - loss: 0.1207 - accuracy: 0.9669 - val_loss: 0.1377 - val_accu-
racy: 0.9697
Epoch 64/100
106/106 [=====] - 0s 5ms/step - loss: 0.1203 - accuracy: 0.9706 - val_loss: 0.1385 - val_accu-
racy: 0.9659
Epoch 65/100
106/106 [=====] - 1s 5ms/step - loss: 0.1250 - accuracy: 0.9659 - val_loss: 0.1363 - val_accu-
racy: 0.9659

```

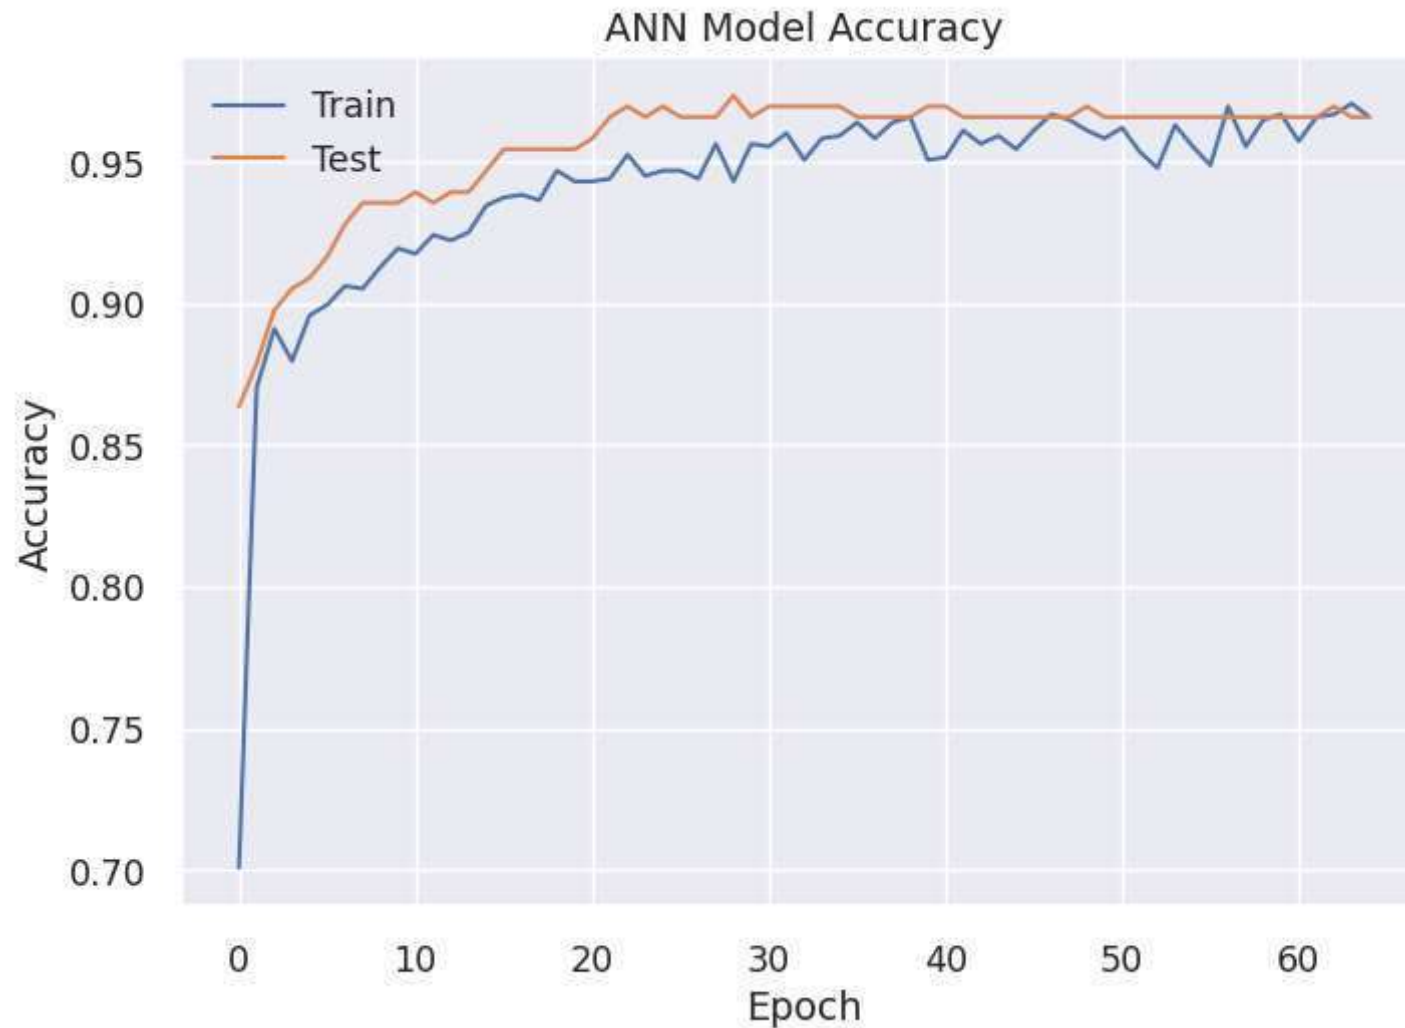
```

In [ ]: # Cell 4: Plot Accuracy History
plt.plot(model_training.history['accuracy'])
plt.plot(model_training.history['val_accuracy'])
plt.title('ANN Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')

```




```
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



```
In [ ]: # Cell 5: Plot Loss History
plt.plot(model_training.history['loss'])
plt.plot(model_training.history['val_loss'])
plt.title('ANN Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
```

```
plt.legend(['Train', 'Test'], loc='upper left')  
plt.show()
```



```
In [ ]: # Cell 6: Predictions and Evaluation  
y_pred_ann = ann.predict(X_test)  
y_pred_ann = (y_pred_ann > 0.5)  
  
acc_ann = accuracy_score(y_test, y_pred_ann)  
print('... ANN Accuracy: {:.2f}%'.format(acc_ann * 100))
```



```
print(classification_report(y_test, y_pred_ann))
conf_matrix_ann = confusion_matrix(y_pred_ann, y_test)
print(conf_matrix_ann)
```

18/18 [=====] - 0s 3ms/step


∴ ANN Accuracy: 95.76%


	precision	recall	f1-score	support
0	0.99	0.91	0.95	252
1	0.93	0.99	0.96	314
accuracy			0.96	566
macro avg	0.96	0.95	0.96	566
weighted avg	0.96	0.96	0.96	566

```
[[230  2]
 [ 22 312]]
```

In []:

```
# Cell 7: Plot Data Distribution
X, y = make_circles(n_samples=150, noise=0.1, random_state=1)
for i in range(2):
    samples_ix = where(y == i)
    plt.scatter(X[samples_ix, 0], X[samples_ix, 1], label=str(i))
plt.legend()
plt.show()
```





Model Comparison 🧐

👉 After implementing 4 models, this section will compare **machine learning models**.

```
In [ ]: # --- Compare All Modules Accuracy ---
compare = pd.DataFrame({
    'Model': ['Logistic Regression', 'Support Vector Machine',
             'Gaussian Naive Bayes', 'Random Forest', 'ANN'],
    'Accuracy': [acc_lr*100, SVMAcc*100, GNBAcc*100, RFAcc*100, acc_ann*100]
})

# --- Display Accuracy Comparison Table ---
compare_sorted = compare.sort_values(by='Accuracy', ascending=False).style.background_gradient(cmap='PuRd').hide_index
compare_sorted
```

```
Out[ ]:
```

Model	Accuracy
ANN	95.759717
Random Forest	95.406360
Support Vector Machine	89.575972
Logistic Regression	85.865724